

DataAnalysisLearning-Data Cleaning and Preparation

July 5, 2019

```
In [1]: # -*- coding: utf-8 -*-
        # @author: tongzi
        # @description: Data Cleanning and Preparation
        # @created date: 2019/07/04
        # @license
```

```
In [2]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        %matplotlib inline
```

```
C:\Anaconda3\lib\importlib\_bootstrap.py:219: RuntimeWarning: numpy.ufunc size changed, may in
    return f(*args, **kws)
C:\Anaconda3\lib\importlib\_bootstrap.py:219: RuntimeWarning: numpy.ufunc size changed, may in
    return f(*args, **kws)
C:\Anaconda3\lib\importlib\_bootstrap.py:219: RuntimeWarning: numpy.ufunc size changed, may in
    return f(*args, **kws)
C:\Anaconda3\lib\importlib\_bootstrap.py:219: RuntimeWarning: numpy.ufunc size changed, may in
    return f(*args, **kws)
C:\Anaconda3\lib\importlib\_bootstrap.py:219: RuntimeWarning: numpy.ufunc size changed, may in
    return f(*args, **kws)
```

In this chapter, we will discuss tools for missing data, duplicate data, string manipulation, and some other analytical data transformations.

0.0.1 Handling Missing Data

Missing data occurs commonly in many data analysis applications. One of the goal of pandas is to make working with missing data as painless as possible. For example, all of the descriptive statistic on pandas objects exclude missing data by default.

For numeric data, pandas uses the float-point value **NaN** (Not a Number) to represent missing data. We call this a sentinel value that can be easily detected:

```
In [3]: string_data = pd.Series(['aardvark', 'artichoke', np.nan, 'avocado'])
```

```
In [4]: string_data
```

```
Out [4]: 0    aardvark
         1    artichoke
         2         NaN
         3    avocado
         dtype: object
```

```
In [5]: string_data.isnull()
```

```
Out [5]: 0    False
         1    False
         2     True
         3    False
         dtype: bool
```

Below is a table that lists some functions related to missing data handling:

0.0.2 Filtering Out Missing Data

```
In [6]: from numpy import nan as NA
```

```
In [7]: data = pd.Series([1, NA, 3.5, 7])
```

```
In [8]: data.dropna()
```

```
Out [8]: 0    1.0
         2    3.5
         3    7.0
         dtype: float64
```

As we can see above, the *dropna()* method returns the Series with only the non-null data and index values. This is equivalent to:

```
In [9]: data[data.notnull()]
```

```
Out [9]: 0    1.0
         2    3.5
         3    7.0
         dtype: float64
```

With DataFrame, things are a bit more different, we may want to drop rows or columns that are all NA or only those containing any NAs. *dropna()* method by default drops any row containing a missing value:

```
In [10]: data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA],
        ....: [NA, NA, NA], [NA, 6.5, 3.]])
```

```
In [11]: data
```

```
Out [11]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

```
In [12]: data.dropna()
```

```
Out [12]:
```

	0	1	2
0	1.0	6.5	3.0

DataFrame*dropna()*

```
In [15]: #
         data.dropna(axis='columns')
```

```
Out [15]: Empty DataFrame
         Columns: []
         Index: [0, 1, 2, 3]
```

```
In [16]: data
```

```
Out [16]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

Passing *how='all'*, will only drop rows that are all NA:

```
In [17]: data.dropna(how='all')
```

```
Out [17]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
3	NaN	6.5	3.0

```
In [18]: #
         data[4] = NA
```

```
In [19]: data
```

```
Out [19]:
```

	0	1	2	4
0	1.0	6.5	3.0	NaN
1	1.0	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	6.5	3.0	NaN

```
In [20]: data.dropna(axis='columns', how='all')
```

```
Out [20]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

DataFrame*thresh*

```
In [21]: df = pd.DataFrame(np.random.randn(7, 3))
```

```
In [22]: df
```

```
Out [22]:
```

	0	1	2
0	-0.377788	-0.099363	0.773704
1	-0.807716	-0.697565	0.309078
2	-1.002575	-0.002337	0.970096
3	1.768859	-0.646667	0.167954
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

```
In [23]: # 41NaN
df.iloc[:4, 1] = NA
```

```
In [26]: # 22NaN
df.iloc[:2, 2] = NA
```

```
In [27]: df
```

```
Out [27]:
```

	0	1	2
0	-0.377788	NaN	NaN
1	-0.807716	NaN	NaN
2	-1.002575	NaN	0.970096
3	1.768859	NaN	0.167954
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

```
In [28]: # NaN
df.dropna()
```

```
Out [28]:
```

	0	1	2
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

```
In [29]: # thresh
#
df.dropna(thresh=2)
```

```
Out [29]:
```

	0	1	2
2	-1.002575	NaN	0.970096
3	1.768859	NaN	0.167954
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

0.0.3 Filling In Missing Data

Rather than filtering out missing data (potentially discarding other data along with it), we may want to fill in the 'hole' in any number of ways. For most purposes, the *fillna()* method is the workhorse function to use:

```
In [30]: # 0
df.fillna(0)
```

```
Out [30]:
```

	0	1	2
0	-0.377788	0.000000	0.000000
1	-0.807716	0.000000	0.000000
2	-1.002575	0.000000	0.970096
3	1.768859	0.000000	0.167954
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

```
In [32]: #
# 1102200
df.fillna({1:10, 2:200})
```

```
Out [32]:
```

	0	1	2
0	-0.377788	10.000000	200.000000
1	-0.807716	10.000000	200.000000
2	-1.002575	10.000000	0.970096
3	1.768859	10.000000	0.167954
4	0.687819	3.742656	-0.482712
5	-0.035491	0.267234	1.161451
6	-0.654256	-0.612303	2.564999

fillna() returns a new object, but we can modify the existing object in-place by passing the argument *inplace*:

```
In [33]: df.fillna(0, inplace=True)
```

```
In [34]: df
```

```
Out [34]:
```

	0	1	2
0	-0.377788	0.000000	0.000000
1	-0.807716	0.000000	0.000000
2	-1.002575	0.000000	0.970096

```

3  1.768859  0.000000  0.167954
4  0.687819  3.742656 -0.482712
5 -0.035491  0.267234  1.161451
6 -0.654256 -0.612303  2.564999

```

The same interpolations available for reindexing can be used for `fillna()`: `>fillna()`

```
In [35]: df = pd.DataFrame(np.random.randn(6, 3))
```

```
In [36]: # implicate indexing
df.iloc[2:, 1] = NA # 21NaN
```

```
In [39]: # implicate indexing
df.iloc[4:, 2] = NA # 42NaN
```

```
In [40]: df
```

```
Out [40]:
```

	0	1	2
0	-0.018569	0.250788	0.468610
1	-1.147574	1.921718	1.831329
2	-1.351726	NaN	-0.247973
3	0.824898	NaN	0.608928
4	-0.540984	NaN	NaN
5	-0.748586	NaN	NaN

```
In [41]: # 'ffill'forward fill,
df.fillna(method='ffill')
```

```
Out [41]:
```

	0	1	2
0	-0.018569	0.250788	0.468610
1	-1.147574	1.921718	1.831329
2	-1.351726	1.921718	-0.247973
3	0.824898	1.921718	0.608928
4	-0.540984	1.921718	0.608928
5	-0.748586	1.921718	0.608928

```
In [42]: # 'ffill'forward fill,,
# 2
df.fillna(method='ffill', limit=2)
```

```
Out [42]:
```

	0	1	2
0	-0.018569	0.250788	0.468610
1	-1.147574	1.921718	1.831329
2	-1.351726	1.921718	-0.247973
3	0.824898	1.921718	0.608928
4	-0.540984	NaN	0.608928
5	-0.748586	NaN	0.608928

With the `fillna()`, we can do lots of things with a little creativity. For example, we might pass the mean or median value of a Series:

```
In [43]: data = pd.Series([1., NA, 3.5, NA, 7])
```

```
In [44]: data
```

```
Out[44]: 0    1.0  
         1    NaN  
         2    3.5  
         3    NaN  
         4    7.0  
         dtype: float64
```

```
In [45]: #(1.0 + 3.5 + 7.0) / 3  
         data.fillna(data.mean())
```

```
Out[45]: 0    1.000000  
         1    3.833333  
         2    3.500000  
         3    3.833333  
         4    7.000000  
         dtype: float64
```

```
In [48]: (1.0 + 3.5 + 7.0) / 3
```

```
Out[48]: 3.8333333333333335
```

fillna()

0.0.4 Data Transformation

So far in this chapter, we've been concerned with rearranging data. Filtering, cleaning and other transformations are another class of important operations.

Removing Duplicates Duplicate row may be found in a DataFrame for any number of reasons. For example:

```
In [49]: data = pd.DataFrame({'k1': ['one', 'two'] * 3 + ['two'],  
                             ....: 'k2': [1, 1, 2, 3, 3, 4, 4]})
```

```
In [50]: data
```

```
Out[50]:   k1  k2  
0  one   1  
1  two   1  
2  one   2  
3  two   3  
4  one   3  
5  two   4  
6  two   4
```

The DataFrame method *duplicated()* returns a Series indicating whether each row is a duplicate (has been observed in a previous row) or not:

```
In [51]: data.duplicated()
```

```
Out[51]: 0    False
         1    False
         2    False
         3    False
         4    False
         5    False
         6     True
         dtype: bool
```

Relatively, the DataFrame *drop_duplicates()* returns a DataFrame where the duplicated array is False:

```
In [52]: data.drop_duplicates()
```

```
Out[52]:   k1  k2
0  one   1
1  two   1
2  one   2
3  two   3
4  one   3
5  two   4
```

```
In [53]: data['v1'] = range(7)
```

```
In [54]: data
```

```
Out[54]:   k1  k2  v1
0  one   1   0
1  two   1   1
2  one   2   2
3  two   3   3
4  one   3   4
5  two   4   5
6  two   4   6
```

```
In [56]: # k1
         data.drop_duplicates(['k1'])
```

```
Out[56]:   k1  k2  v1
0  one   1   0
1  two   1   1
```

*duplicated()**drop_duplicates()**keep='last'*


```
In [57]: data.drop_duplicates(['k1', 'k2'], keep='last')
```

```
Out[57]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
6	two	4	6

Transformationg Data Using a Function or Mapping SeriesDataFrame

```
In [58]: data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',  
....: 'Pastrami', 'corned beef', 'Bacon',  
....: 'pastrami', 'honey ham', 'nova lox'],  
....: 'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
```

```
In [59]: data
```

```
Out[59]:
```

	food	ounces
0	bacon	4.0
1	pulled pork	3.0
2	bacon	12.0
3	Pastrami	6.0
4	corned beef	7.5
5	Bacon	8.0
6	pastrami	3.0
7	honey ham	5.0
8	nova lox	6.0

>Suppose we want to add a column indicating the type of the animal that each food came from.

```
In [60]: meat_to_animal = {  
    'bacon': 'pig',  
    'pulled pork': 'pig',  
    'pastrami': 'cow',  
    'corned beef': 'cow',  
    'honey ham': 'pig',  
    'nova lox': 'salmon'  
}
```

The `map()` method on a Series accepts a function or dict-like containing a mapping, but here we have a problem in that some of the meats are capitalized and others are not. Thus, we need to convert each value to lowercase using the `Series.str.lower()` method:

```
In [61]: lowercased = data['food'].str.lower()
```

```
In [62]: lowercased
```

```
Out[62]: 0      bacon
1  pulled pork
2      bacon
3    pastrami
4  corned beef
5      bacon
6    pastrami
7  honey ham
8    nova lox
Name: food, dtype: object
```

```
In [63]: # animal
data['animal'] = lowercased.map(meat_to_animal)
```

```
In [64]: data
```

```
Out[64]:
```

	food	ounces	animal
0	bacon	4.0	pig
1	pulled pork	3.0	pig
2	bacon	12.0	pig
3	Pastrami	6.0	cow
4	corned beef	7.5	cow
5	Bacon	8.0	pig
6	pastrami	3.0	cow
7	honey ham	5.0	pig
8	nova lox	6.0	salmon

We could also pass a function that does the same work:

```
In [67]: data['food'].map(lambda x: meat_to_animal[x.lower()])
```

```
Out[67]: 0      pig
1      pig
2      pig
3      cow
4      cow
5      pig
6      cow
7      pig
8  salmon
Name: food, dtype: object
```

Using *map()* method is a convenient way to perform element-wise transformation and other related data cleaning-related operations.

Replacing Values *fillna()**map()**replace()* Let's consider this Series:

```
In [69]: data = pd.Series([1., -999., 2., -999., -1000., 3.])
```

```
In [70]: data
```

```
Out[70]: 0      1.0  
         1    -999.0  
         2      2.0  
         3    -999.0  
         4   -1000.0  
         5      3.0  
         dtype: float64
```

-999.0pandasreplace(): inplace=True.

```
In [71]: data.replace(-999, np.nan)
```

```
Out[71]: 0      1.0  
         1     NaN  
         2      2.0  
         3     NaN  
         4   -1000.0  
         5      3.0  
         dtype: float64
```

If we want to replace multiple values at once, we instead pass a list and then the substitute value():

```
In [72]: data.replace([-999, -1000], np.nan)
```

```
Out[72]: 0      1.0  
         1     NaN  
         2      2.0  
         3     NaN  
         4     NaN  
         5      3.0  
         dtype: float64
```

To use a different replacement for each value, pass a list of substitute:

```
In [74]: # -999np.nan  
         # -10000  
         data.replace([-999, -1000], [np.nan, 0])
```

```
Out[74]: 0      1.0  
         1     NaN  
         2      2.0  
         3     NaN  
         4      0.0  
         5      3.0  
         dtype: float64
```

The argument passed can also be a dict:

```
In [75]: data.replace({-999: np.nan, -1000:0})
```

```
Out[75]: 0    1.0
         1    NaN
         2    2.0
         3    NaN
         4    0.0
         5    3.0
         dtype: float64
```

The `data.replace()` is distinct from `data.str.replace()`, which performs string substitute element-wise.

Renaming Axis Indexes SeriesDataFrame

```
In [76]: data = pd.DataFrame(np.arange(12).reshape((3, 4)),
    ....: index=['Ohio', 'Colorado', 'New York'],
    ....: columns=['one', 'two', 'three', 'four'])
```

```
In [77]: data
```

```
Out[77]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
New York	8	9	10	11

Seriesmap()

```
In [79]: #
         data.index.map(lambda x: x.upper())
```

```
Out[79]: Index(['OHIO', 'COLORADO', 'NEW YORK'], dtype='object')
```

```
In [81]: data.index = data.index.map(lambda x: x.upper())
```

```
In [82]: data
```

```
Out[82]:
```

	one	two	three	four
OHIO	0	1	2	3
COLORADO	4	5	6	7
NEW YORK	8	9	10	11

If we want to create a transformed version of dataset without modifying the original, a useful method is `rename()`:

```
In [85]: #
         #
         data.rename(index=str.title, columns=str.upper)
```

```
Out [85]:
```

	ONE	TWO	THREE	FOUR
Ohio	0	1	2	3
Colorado	4	5	6	7
New York	8	9	10	11

```
rename()
```

```
In [88]: data
```

```
Out [88]:
```

	one	two	three	four
OHIO	0	1	2	3
COLORADO	4	5	6	7
NEW YORK	8	9	10	11

```
In [89]: data.rename(index={'NEW YORK': 'Nanning'},
                      columns={'four': 4})
```

```
Out [89]:
```

	one	two	three	4
OHIO	0	1	2	3
COLORADO	4	5	6	7
Nanning	8	9	10	11

`rename()` method saves us from the chore () of copying the DataFrame manually and assigning to its index and columns attributes. By passing the argument *inplace*, we can modify a dataset in-place:

```
In [90]: data.rename(index={'OHIO': 'Liuzhou'}, inplace=True)
```

```
In [91]: data
```

```
Out [91]:
```

	one	two	three	four
Liuzhou	0	1	2	3
COLORADO	4	5	6	7
NEW YORK	8	9	10	11

Discretization and Binning

Continuous data is discretized or otherwise separated into “bins” for analysis. Suppose we have data about a group of people in a study, and we want to group them into discrete age buckets ():

```
In [92]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

Let’s divide these into bins of 18 to 25, 26 to 35, 36 to 60, and finally 61 and older. To do so, we have to use `cut()`, a method in pandas:

```
In [93]: bins = [18, 25, 35, 60, 100]
```

```
In [94]: cats = pd.cut(ages, bins)
```

```
In [95]: cats
```

```
Out[95]: [(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (35, 60)]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

The object pandas returns is a special Categorical object. The object's *codes* attribute contains a categories array specifying the distinct category names along with a labeling for the *ages* data:

```
In [97]: # 40,1,2,3,4
cats.codes
```

```
Out[97]: array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)
```

```
In [99]: #
cats.categories
```

```
Out[99]: IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]],
                        closed='right',
                        dtype='interval[int64]')
```

pandas.value_counts()

```
In [100]: pd.value_counts(cats)
```

```
Out[100]: (18, 25]      5
          (35, 60]      3
          (25, 35]      3
          (60, 100]     1
          dtype: int64
```

Consistent with mathematical notation for intervals (), a parenthesis means that the side is open, while the square bracket mean it is closed (inclusive). We can change which side is closed by passing *right=False*:

```
In [101]: pd.cut(ages, [18, 26, 36, 61, 100], right=False)
```

```
Out[101]: [[18, 26), [18, 26), [18, 26), [26, 36), [18, 26), ..., [26, 36), [61, 100), [36, 61)]
Length: 12
Categories (4, interval[int64]): [[18, 26) < [26, 36) < [36, 61) < [61, 100))
```

We can also pass our own bin names by passing a list or an array to the *labels* option:

```
In [102]: group_names = ["Youth", "YoungAdult", "MiddleAged", "Senior"]
```

```
In [103]: pd.cut(ages, bins, labels=group_names)
```

```
Out[103]: [Youth, Youth, Youth, YoungAdult, Youth, ..., YoungAdult, Senior, MiddleAged, MiddleAged]
Length: 12
Categories (4, object): [Youth < YoungAdult < MiddleAged < Senior]
```

If we pass an integer number of bins to the *cut()* method instead of explicit edges, it will compute the equal-length bins based on the minimum and maximum values in the data. Consider the case of some uniformly distributed data chopped into fourths:

```
In [104]: data = np.random.rand(20)
```

```
In [106]: # data 4 bins  
# 2  
pd.cut(data, 4, precision=2)
```

```
Out[106]: [(0.0023, 0.25], (0.49, 0.74], (0.0023, 0.25], (0.0023, 0.25], (0.0023, 0.25], ...  
Length: 20  
Categories (4, interval[float64]): [(0.0023, 0.25] < (0.25, 0.49] < (0.49, 0.74] < (0.74, 1.0]
```

The *precision=2* option limits the decimal precision to two digits.

A closely related function, *qcut()*, bins the data on sample quantiles. Depending on the distribution of the data, using *cut()* method will not usually result in each bin having same number of data point. Since *qcut()* uses sample quantiles instead, by definition you will obtain roughly equal-size bins:

```
In [107]: data = np.random.randn(1000) # normal distribution
```

```
In [116]: cats = pd.qcut(data, 4) # cut into quartile
```

```
In [117]: cats
```

```
Out[117]: [(-0.00142, 0.631], (-3.009, -0.611], (-0.611, -0.00142], (-0.611, -0.00142], (-0.00142, 0.631]  
Length: 1000  
Categories (4, interval[float64]): [(-3.009, -0.611] < (-0.611, -0.00142] < (-0.00142, 0.631] < (0.631, 3.77]
```

```
In [118]: pd.value_counts(cats)
```

```
Out[118]: (0.631, 3.77]          250  
(-0.00142, 0.631]          250  
(-0.611, -0.00142]          250  
(-3.009, -0.611]           250  
dtype: int64
```

Similar to *cut()*, we can pass our own quantiles (numbers between 0 and 1, inclusive):

```
In [119]: pd.qcut(data, [0, 0.1, 0.5, 0.6, 0.9, 1.])
```

```
Out[119]: [(-0.00142, 0.215], (-3.009, -1.253], (-1.253, -0.00142], (-1.253, -0.00142], (-0.00142, 0.215]  
Length: 1000  
Categories (5, interval[float64]): [(-3.009, -1.253] < (-1.253, -0.00142] < (-0.00142, 0.215] < (0.215, 0.631] < (0.631, 3.77]
```

Detecting and Filtering Outliers ()

```
In [120]: data = pd.DataFrame(np.random.randn(1000, 4))
```

```
In [121]: data.describe()
```

```
Out [121]:
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.062161	0.010908	-0.012074	0.010671
std	0.994697	1.007757	0.964354	0.999827
min	-3.299686	-2.952604	-3.579078	-3.577651
25%	-0.620926	-0.645775	-0.700070	-0.668896
50%	0.014344	-0.017357	0.012803	0.038414
75%	0.764888	0.677508	0.626660	0.709005
max	3.192125	3.342354	3.205287	2.542400

Suppose we want to find values in one of the columns exceeding 3 in absolute value:

```
In [122]: col = data[2]
```

```
In [123]: col[np.abs(col) > 3]
```

```
Out [123]: 189    3.109010
389   -3.579078
824    3.205287
Name: 2, dtype: float64
```

3DataFrame.any()

```
In [124]: data[(np.abs(data) > 3).any(1)]
```

```
Out [124]:
```

	0	1	2	3
189	-0.513485	0.464559	3.109010	0.958465
389	-0.377858	0.341131	-3.579078	-0.906682
450	0.516619	0.099933	-0.835361	-3.163249
599	-2.139750	1.593949	1.175514	-3.435382
638	-1.383712	3.318922	-0.770849	-2.257575
668	-0.805004	0.233830	0.284846	-3.577651
735	-3.299686	0.157850	1.325712	-0.321169
824	0.266576	2.287888	3.205287	1.195037
844	-0.756320	3.342354	-2.250097	-1.664915
852	3.192125	-0.897974	0.696980	0.707895
953	1.806269	0.304860	-0.447591	-3.053805

[-3,3]3

```
In [125]: data[np.abs(data) > 3] = np.sign(data) * 3
```

```
In [126]: data.describe()
```

```
Out [126]:
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.062269	0.010246	-0.011809	0.011901
std	0.993141	1.005683	0.961369	0.995834
min	-3.000000	-2.952604	-3.000000	-3.000000
25%	-0.620926	-0.645775	-0.700070	-0.668896
50%	0.014344	-0.017357	0.012803	0.038414
75%	0.764888	0.677508	0.626660	0.709005
max	3.000000	3.000000	3.000000	2.542400


```
In [130]: test = np.random.randint(-10, 10, size=(3,3))
```

```
In [136]: test = pd.DataFrame(test)
```

```
In [137]: test
```

```
Out[137]:
```

	0	1	2
0	-8	-2	-9
1	-4	7	-3
2	0	8	-8

```
In [139]: test[np.abs(test) > 3] = np.sign(test) * 3
```

```
In [140]: test
```

```
Out[140]:
```

	0	1	2
0	-3	-2	-3
1	-3	3	-3
2	0	3	-3

```
In [141]: np.sign(test) * 3
```

```
Out[141]:
```

	0	1	2
0	-3	-3	-3
1	-3	3	-3
2	0	3	-3

np.sign(data)
data-11(shape)
data

Permutation and Random Sampling Permuting (randomly reordering) a Series or the rows of a DataFrame is easy to do using the *numpy.random.permutation()*. Calling *permutation()* with the length of the axis you want to permute produces an array of integers indicating the new ordering:

```
In [142]: df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))
```

```
In [143]: df
```

```
Out[143]:
```

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

```
In [146]: sampler = np.random.permutation(5)
```

```
In [147]: sampler
```

```
Out[147]: array([3, 2, 0, 1, 4])
```

The array can be used in `iloc`-based indexing or the equivalent function `take()` function:

```
In [148]: df.iloc[sampler]
```

```
Out[148]:
```

	0	1	2	3
3	12	13	14	15
2	8	9	10	11
0	0	1	2	3
1	4	5	6	7
4	16	17	18	19

```
In [149]: df.take(sampler)
```

```
Out[149]:
```

	0	1	2	3
3	12	13	14	15
2	8	9	10	11
0	0	1	2	3
1	4	5	6	7
4	16	17	18	19

To select a random subset without replacement, we can use the `sample()` method on Series and DataFrame:

```
In [152]: #
          df.sample(n=3)
```

```
Out[152]:
```

	0	1	2	3
4	16	17	18	19
2	8	9	10	11
3	12	13	14	15

```
In [156]: df.sample(n=2, axis=1)
```

```
Out[156]:
```

	2	0
0	2	0
1	6	4
2	10	8
3	14	12
4	18	16

To generate a sample with replacement (to allow repeat values), pass the argument `replace=True` to `sample`:

```
In [157]: choices = pd.Series([5, 7, -1, 6, 4])
```

```
In [159]: draws = choices.sample(n=10, replace=True)
```

```
In [160]: draws
```

```
Out[160]: 0    5
          3    6
          2   -1
          4    4
          1    7
          3    6
          2   -1
          4    4
          0    5
          3    6
          dtype: int64
```

Computing Indicator/Dummy Variables DataFrame(k)k01pandas.get_dummies()

```
In [161]: df = pd.DataFrame({'key': ['b', 'b', 'a', 'c', 'a', 'b'],
        .....: 'data1': range(6)})
```

```
In [162]: df
```

```
Out[162]:   key  data1
0    b      0
1    b      1
2    a      2
3    c      3
4    a      4
5    b      5
```

```
In [163]: pd.get_dummies(df['key'])
```

```
Out[163]:    a  b  c
0  0  1  0
1  0  1  0
2  1  0  0
3  0  0  1
4  1  0  0
5  0  1  0
```

```
dfkeya, b, cpandas.get_dummies()010aa(df)
,pandas.get_dummies()prefix
```

```
In [164]: dummies = pd.get_dummies(df['key'], prefix='key_')
```

```
In [165]: dummies
```

```
Out[165]:   key__a  key__b  key__c
0         0         1         0
1         0         1         0
2         1         0         0
3         0         0         1
4         1         0         0
5         0         1         0
```

```

df['data1']

In [174]: df['data1'] # Series

Out[174]: 0      0
          1      1
          2      2
          3      3
          4      4
          5      5
          Name: data1, dtype: int64

In [175]: df[['data1']] # DataFrame

Out[175]:   data1
0         0
1         1
2         2
3         3
4         4
5         5

In [176]: df[['data1']].join(dummies) # DataFrame

Out[176]:   data1  key__a  key__b  key__c
0         0         0         1         0
1         1         0         1         0
2         2         1         0         0
3         3         0         0         1
4         4         1         0         0
5         5         0         1         0

In [179]: #
          data.rename(columns={0:'one', 1:'two', 2:'three', 3:'four'}, inplace=True)

In [185]: # 'one' 'two' DataFrame
          dd = data[['one', 'four']]

In [186]: # 'one' Series
          data['one']

Out[186]: 0      -1.486255
          1      -0.176648
          2      -1.150700
          3      -0.769186
          4      -0.882545
          5       0.043969
          6      -0.941941
          7       0.919140

```

8	-0.591679
9	0.603175
10	-0.005484
11	0.294564
12	-1.446958
13	-0.513096
14	0.037718
15	-0.107953
16	0.613648
17	0.156419
18	-0.725193
19	-0.385183
20	0.641327
21	-1.110199
22	-0.237171
23	0.072884
24	0.460074
25	-0.541733
26	0.057325
27	1.426806
28	-1.200798
29	-0.142353
	...
970	0.445650
971	-0.055439
972	1.682989
973	0.123819
974	2.127047
975	0.710947
976	0.279677
977	-0.419060
978	2.621613
979	0.270127
980	0.776597
981	0.041169
982	-0.337197
983	0.828886
984	-0.569314
985	0.886491
986	1.532338
987	1.209169
988	0.263185
989	0.709655
990	0.224568
991	-0.319106
992	1.229200
993	1.168087
994	2.138031

```

995    0.211693
996    1.655879
997    0.591293
998   -1.533716
999    0.300527
Name: one, Length: 1000, dtype: float64

```

If a row in a DataFrame belongs to multiple categories, things are a bit more complicated.

```

In [189]: # 'one' 'four' data.columns
          data.columns.get_indexer(['one', 'four'])

```

```

Out[189]: array([0, 3], dtype=int64)

```

A useful recipe for staticstical applications is to combine `get_dummies()` with a discretization function like `cut()`:

```

In [190]: rng = np.random.RandomState(12345) #

```

```

In [191]: values = rng.rand(10)

```

```

In [192]: values

```

```

Out[192]: array([0.92961609, 0.31637555, 0.18391881, 0.20456028, 0.56772503,
                 0.5955447 , 0.96451452, 0.6531771 , 0.74890664, 0.65356987])

```

```

In [193]: #
          #
          bins = [0, 0.2, 0.4, 0.6, 0.8, 1.0]

```

```

In [194]: cats = pd.cut(values, bins)

```

```

In [195]: cats

```

```

Out[195]: [(0.8, 1.0], (0.2, 0.4], (0.0, 0.2], (0.2, 0.4], (0.4, 0.6], (0.4, 0.6], (0.8, 1.0],
Categories (5, interval[float64]): [(0.0, 0.2] < (0.2, 0.4] < (0.4, 0.6] < (0.6, 0.8]

```

```

In [196]: pd.get_dummies(cats)

```

```

Out[196]:
```

	(0.0, 0.2]	(0.2, 0.4]	(0.4, 0.6]	(0.6, 0.8]	(0.8, 1.0]
0	0	0	0	0	1
1	0	1	0	0	0
2	1	0	0	0	0
3	0	1	0	0	0
4	0	0	1	0	0
5	0	0	1	0	0
6	0	0	0	0	1
7	0	0	0	1	0
8	0	0	0	1	0
9	0	0	0	1	0

We create a random instance with seed=12345 to make the result deterministic.

String Manipulation

Python has long been a popular raw data manipulation language in part due to its ease of use for string and text processing. Many text operations are made simple with the string object's built-in methods. For more complex pattern matching and text manipulation, regular expressions may be needed. *pandas adds to the mix by enabling us to apply string and regular expressions concisely on whole arrays of data, additionally handling the annoyance of missing data.

String Object Methods In many string munging and scripting applications, built-in methods are sufficient. For example, a comma-separated string can be broken into pieces with *split()*:

```
In [197]: val = 'a,b,    guido'
```

```
In [198]: val.split(',')
```

```
Out[198]: ['a', 'b', '    guido']
```

split() often combined with *strip()* to trim whitespace (including line breaks):

```
In [199]: pieces = [x.strip() for x in val.split(',')] 
```

```
In [200]: pieces
```

```
Out[200]: ['a', 'b', 'guido']
```

These strings could be concatenated together with a two-colons delimiter using addition:

```
In [201]: first, second, third = pieces
```

```
In [202]: first + "::" + second + "::" + third
```

```
Out[202]: 'a::b::guido'
```

But this isn't a practical general method. A faster and more Pythonic is to pass a tuple or list to the *join()* method on the string "::":

```
In [203]: "::".join(pieces)
```

```
Out[203]: 'a::b::guido'
```

*Pythoninindex()**find()*

```
In [204]: 'guido' in val
```

```
Out[204]: True
```

```
In [205]: val.index(',')
```

```
Out[205]: 1
```

```
In [206]: val.find("::")
```

```
Out[206]: -1
```

```
index()find()index()find()-1
```

Relatedly, `count()` method returns the number of occurrences of a particular substring:

```
In [207]: val.count(',')
```

```
Out[207]: 2
```

`replace()` method will substitute () the occurrences of one pattern for another. It is common to use delete patterns, too, by passing an empty string:

```
In [208]: # val', '::'  
val.replace(',', '::')
```

```
Out[208]: 'a::b::    guido'
```

```
In [209]: # val''  
val.replace(',', '')
```

```
Out[209]: 'ab    guido'
```

Python built-in string methods:

Regular Expressions Regular expressions provide a flexible way to search and match (often more complex) string patterns in text. Python's built-in module `re` is responsible for applying regular expressions to strings.

The `re` module falls into three categories: pattern matching, substitution, and splitting. Naturally these are all related; a regex (short for regular expression) describes a pattern to locate in the text, which can be used for many purposes. Let's look at an example:

Suppose we want to split a string with a variable number of whitespaces (tabs, spaces and newlines). The regex describing one or more whitespace characters is `+`:

```
In [210]: import re
```

```
In [211]: text = 'foo    bar\t baz \tqux'
```

```
In [212]: re.split('\s+', text)
```

```
Out[212]: ['foo', 'bar', 'baz', 'qux']
```

```
re.split('+', text)split()textre.compile()regex:
```

```
In [213]: regex = re.compile('\s+')
```

```
In [214]: regex.split(text)
```



```
Out [214]: ['foo', 'bar', 'baz', 'qux']
```

```
re.findall()
```

```
In [215]: regex.findall(text)
```

```
Out [215]: [' ', '\t ', ' \t']
```

Creating a regex object with *re.compile()* method is highly recommended if want to apply the same expression to many strings; doing so will save many CPU cycles.

```
match()search()findall()findall()search()match()
```

Let's consider a block of text and a regular expression capable of identifying most email addresses:

```
In [216]: text = """Dave dave@google.com
Steve steve@gmail.com
Rob rob@gmail.com
Ryan ryan@yahoo.com
"""
```

```
In [223]: pattern = r'[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,4}'
```

```
In [224]: #
# re.IGNORECASE makes the regex case-insensitive
regex = re.compile(pattern, flags=re.IGNORECASE)
```

```
findall():
```

```
In [225]: regex.findall(text)
```

```
Out [225]: ['dave@google.com', 'steve@gmail.com', 'rob@gmail.com', 'ryan@yahoo.com']
```

```
search()match:
```

```
In [226]: m = regex.search(text)
```

```
In [227]: m
```

```
Out [227]: <re.Match object; span=(5, 20), match='dave@google.com'>
```

```
In [228]: text[m.start():m.end()]
```

```
Out [228]: 'dave@google.com'
```

```
In [230]: print(regex.match(text))
```

```
None
```

regex.match() returns None, as it only matches if the pattern occurs at the beginning of the string.

Relatedly, *sub()* method will return a new string with occurrences of the string replaced by the new string:

```
In [231]: print(regex.sub("redacted", text))
```

```
Dave redacted  
Steve redacted  
Rob redacted  
Ryan redacted
```

(parenthesis)

```
In [233]: pattern = r'([A-Z0-9._%+-]+)@([A-Z0-9.-]+)\.([A-Z]{2,4})'
```

```
In [237]: regex = re.compile(pattern, flags=re.IGNORECASE)
```

A match object produced by this modified regex returns a tuple of the pattern components with its *group()* method:

```
In [238]: m = regex.match('tongzi@126.com')
```

```
In [239]: m.groups()
```

```
Out[239]: ('tongzi', '126', 'com')
```

While *findall()* returns a list of tuples when the pattern has groups:

```
In [240]: regex.findall(text)
```

```
Out[240]: [('dave', 'google', 'com'),  
            ('steve', 'gmail', 'com'),  
            ('rob', 'gmail', 'com'),  
            ('ryan', 'yahoo', 'com')]
```

sub() also has access to groups in each match using special symbols like \1 and \2. The symbol \1 corresponds to the first matched group, \2 corresponds to the second, and so forth:

```
In [241]: print(regex.sub(r'User: \1, Domain: \2, Suffix: \3', text))
```

```
Dave User: dave, Domain: google, Suffix: com  
Steve User: steve, Domain: gmail, Suffix: com  
Rob User: rob, Domain: gmail, Suffix: com  
Ryan User: ryan, Domain: yahoo, Suffix: com
```

Table below provides a brief summary:

Vectorized String Functions in pandas Cleaning up a messy dataset for analysis often requires a lot of string munging and regularization. To complicate matters, a column containing strings will sometimes have missing data:

```
In [242]: data = {'Dave': 'dave@google.com', 'Steve': 'steve@gmail.com',  
.....: 'Rob': 'rob@gmail.com', 'Wes': np.nan}
```

```
In [243]: data = pd.Series(data)
```

```
In [244]: data
```

```
Out[244]: Dave      dave@google.com  
         Steve    steve@gmail.com  
         Rob      rob@gmail.com  
         Wes      NaN  
         dtype: object
```

```
In [245]: data.isnull()
```

```
Out[245]: Dave      False  
         Steve    False  
         Rob      False  
         Wes      True  
         dtype: bool
```

We can apply string and regular expression methods that can be applied (passing a lambda or other function) to each value using *data.map()*, but it will fail on NA values. To code with this, Series has array-oriented methods for string operations that skip NA values. These methods are accessed through Series's *str* attribute. For example, we can check whether each email address has 'gmail' in it with *str.contains()*:

```
In [246]: data.str.contains('gmail')
```

```
Out[246]: Dave      False  
         Steve     True  
         Rob       True  
         Wes      NaN  
         dtype: object
```

Regular expressions can be used, too, along with any *re* options like *re.IGNORECASE*:

```
In [248]: pattern
```

```
Out[248]: '([A-Z0-9._%+-]+)@([A-Z0-9.-]+)\.([A-Z]{2,4})'
```

```
In [249]: data.str.findall(pattern, flags=re.IGNORECASE)
```

```
Out[249]: Dave      [(dave, google, com)]  
         Steve    [(steve, gmail, com)]  
         Rob      [(rob, gmail, com)]  
         Wes      NaN  
         dtype: object
```

There are a couple of ways to do vectorized element retrieval. Either use *str.get()* or index into the *str* attribute:

```
In [250]: matches = data.str.match(pattern, flags=re.IGNORECASE)
```

```
In [260]: matches
```

```
Out[260]: Dave      True
          Steve     True
          Rob       True
          Wes       NaN
          dtype: object
```

```
In [255]: data['ceprei'] = 'cepreitest software@ceprei.biz'
```

```
In [256]: data
```

```
Out[256]: Dave                        dave@google.com
          Steve                       steve@gmail.com
          Rob                         rob@gmail.com
          Wes                         NaN
          ceprei  cepreitest software@ceprei.biz
          dtype: object
```

```
In [257]: data.str.match(pattern, flags=re.IGNORECASE)
```

```
Out[257]: Dave      True
          Steve     True
          Rob       True
          Wes       NaN
          ceprei    False
          dtype: object
```

`cepreicepreitestdata.str.match()`False

To access elements in the embedded lists, we can pass an index to either of these functions:

```
In [258]: # matchesSeries
          matches.str.get(1)
```

```
Out[258]: Dave      NaN
          Steve     NaN
          Rob       NaN
          Wes       NaN
          dtype: float64
```

```
In [261]: matches.str[0]
```

```
Out[261]: Dave      NaN
          Steve     NaN
          Rob       NaN
          Wes       NaN
          dtype: float64
```

```
In [262]: data.str[:5]
```

```
Out[262]: Dave      dave@  
         Steve     steve  
         Rob       rob@g  
         Wes       NaN  
         ceprei    cepre  
         dtype: object
```

Partial listing of vectorized string methods: